Chapter 1

Objective of the project

The goal of this project is to understand how training a language model such as GPT2-Xl to convert it to COMET can affect its performance in other language modelling tasks when compared to its 'untrained' form. This will help us identify if this kind of training leads to a loss or a gain in capability for a language model. This is done by comparing the performance of the COMET-GPT2-XL model against the 'naive' GPT2-XL on a variety of common modelling tasks and metrics.